

Recall that we've been discussing beliefs. Here are a couple of beliefs that are worthwhile to check. (examples courtesy Dawson Engler.)

**Redundancy Checking.** 1) Code ought to do something. So, when you have code that doesn't do anything, that's suspicious. Look for identity operations, e.g.

`x = x, 1 * y, x&x, x|x.`

Or, a longer example:

```
1      /* 2.4.5-ac8/net/appletalk/aarp.c */
2      da.s_node = sa.s_node;
3      da.s_net = da.s_net;
```

Also, look for unread writes:

```
1      for (entry=priv->lec_arp_tables[i];
2           entry != NULL; entry=next) {
3          next = entry->next; // never read!
4          ...
5      }
```

Redundancy suggests conceptual confusion.

So far, we've talked about MUST-beliefs; violations are clearly wrong (in some sense). Let's examine MAY beliefs next. For such beliefs, we need more evidence to convict the program.

**Process for verifying MAY beliefs.** We proceed as follows:

1. Record every successful MAY-belief check as "check".
2. Record every unsuccessful belief check as "error".
3. Rank errors based on "check" : "error" ratio.

Most likely errors occur when "check" is large, "error" small.

**Example.** One example of a belief is use-after-free:

```
1      free(p);
2      print(*p);
```

That particular case is a MUST-belief. However, other resources are freed by custom (undocumented) free functions. It's hard to get a list of what is a free function and what isn't. So, let's derive them behaviourally.

**Inferring beliefs: finding custom free functions.** The key idea is: if pointer `p` is not used after calling `foo(p)`, then derive a MAY belief that `foo(p)` frees `p`.

OK, so which functions are free functions? Well, just assume all functions free all arguments:

- emit “check” at every call site;
- emit “error” at every use.

(in reality, filter functions with suggestive names).

Putting that into practice, we might observe:

<code>foo(p)</code>	<code>foo(p)</code>	<code>foo(p)</code>	<code>bar(p)</code>	<code>bar(p)</code>	<code>bar(p)</code>
<code>p = x;</code>	<code>p = x;</code>	<code>p = x;</code>	<code>p = 0;</code>	<code>p=0;</code>	<code>p = x;</code>

We would then rank `bar`’s error first. Plausible results might be: 23 free errors, 11 false positives.

**Inferring beliefs: finding routines that may return NULL.** The situation: we want to know which routines may return NULL. Can we use static analysis to find out?

- sadly, this is difficult to know statically (“`return p->next;`”?) and,
- we get false positives: some functions return NULL under special cases only.

Instead, let’s observe what the programmer does. Again, rank errors based on checks vs non-checks. As a first approximation, assume **all** functions can return NULL.

- if pointer checked before use: emit “check”;
- if pointer used before check: emit “error”.

This time, we might observe:

<code>p = bar(...);</code>	<code>p = bar(...);</code>	<code>p = bar(...);</code>	<code>p = bar(...);</code>
<code>p = x;</code>	<code>if (!p) return;</code>	<code>if (!p) return;</code>	<code>if (!p) return;</code>
	<code>p = x;</code>	<code>p = x;</code>	<code>p = x;</code>

Again, sort errors based on the “check”:“error” ratio.

Plausible results: 152 free errors, 16 false positives.

## General statistical technique

When we write “a(); ... b();”, we mean a MAY-belief that a() is followed by b(). We don’t actually know that this is a valid belief. It’s a hypothesis, and we’ll try it out. Algorithm:

- assume every a–b is a valid pair;
- emit “check” for each path with “a()” and then “b()”;
- emit “error” for each path with “a()” and no “b()”.

(actually, prefilter functions that look paired).

Consider:

foo(p, ...);		foo(p, ...);		foo(p, ...);
bar(p, ...); // check		bar(p, ...); // check		// error: foo, no bar!

This applies to the course project as well.

```
1 void scope1() {
2   A(); B(); C(); D();
3 }
4
5 void scope2() {
6   A(); C(); D();
7 }
8
9 void scope3() {
10  A(); B();
11 }
12
13 void scope4() {
14   B(); D(); scope1();
15 }
16
17 void scope5() {
18   B(); D(); A();
19 }
20
21 void scope6() {
22   B(); D();
23 }
```

“A() and B() must be paired”:  
either A() then B() or B() then A().

**Support** = # times a pair of functions appears together.  
 $\text{support}(\{A,B\})=3$

**Confidence**( $\{A,B\},\{A\}$ ) =  
 $\text{support}(\{A,B\})/\text{support}(\{A\}) = 3/4$

Sample output for support threshold 3, confidence threshold 65% (intra-procedural analysis):

- bug:A in scope2, pair: (A B), support: 3, confidence: 75.00%
- bug:A in scope3, pair: (A D), support: 3, confidence: 75.00%
- bug:B in scope3, pair: (B D), support: 4, confidence: 80.00%
- bug:D in scope2, pair: (B D), support: 4, confidence: 80.00%

The point is to find examples like the one from `cmpci.c` where there’s a `lock_kernel()` call, but, on an exceptional path, no `unlock_kernel()` call.

**Summary: Belief Analysis.** We don’t know what the right spec is. So, look for contradictions.

- MUST-beliefs: contradictions = errors!
- MAY-beliefs: pretend they’re MUST, rank by confidence.

(A key assumption behind this belief analysis technique: most of the code is correct.)

**Further references.** Dawson R. Engler, David Yu Chen, Seth Hallem, Andy Chou and Benjamin Chelf. “Bugs as Deviant Behaviors: A general approach to inferring errors in systems code”. In SOSP ’01.

Dawson R. Engler, Benjamin Chelf, Andy Chou, and Seth Hallem. “Checking system rules using system-specific, programmer-written compiler extensions”. In OSDI ’00 (best paper). [www.stanford.edu/~engler/mc-osdi.pdf](http://www.stanford.edu/~engler/mc-osdi.pdf)

Junfeng Yang, Can Sar and Dawson Engler. “eXplode: a Lightweight, General system for Finding Serious Storage System Errors”. In OSDI’06. [www.stanford.edu/~engler/explode-osdi06.pdf](http://www.stanford.edu/~engler/explode-osdi06.pdf)

## Using Linters

We will also talk about linters in this lecture, based on Jamie Wong’s blog post [jamie-wong.com/2015/02/02/linters-as-invariants/](http://jamie-wong.com/2015/02/02/linters-as-invariants/).

**First there was C.** In statically-typed languages, like C,

```
1 #include <stdio.h>
2
3 int main() {
4     printf("%d\n", num);
5     return 0;
6 }
```

the compiler saves you from yourself. The guaranteed invariant:

“if code compiles, all symbols resolve.”

**Less-nice languages.** OK, so you try to run that in JavaScript and it crashes right away. Invariant?

“if code runs, all symbols resolve?”

But what about this:

```
1 function main(x) {
2     if (x) {
3         console.log("Yay");
4     } else {
5         console.log(num);
6     }
7 }
8
9 main(true);
```

Nope! The above invariant doesn't work.

OK, what about this invariant:

“if code runs without crashing, all symbols referenced in the code path executed resolve?”

Nope!

```
1 function main() {
2   try {
3     console.log(num);
4   } catch (err) {
5     console.log("nothing to see here");
6   }
7 }
8
9 main();
```

So, when you're working in JavaScript and maintaining old code, you always have to deduce:

- is this variable defined?
- is this variable always defined?
- do I need to load a script to define that variable?

We have computers. They're powerful. Why is this the developer's problem?!

```
1 function main(x) {
2   if (x) {
3     console.log("Yay");
4   } else {
5     console.log(num);
6   }
7 }
8
9 main(true);
```

Now:

```
$ nodejs /usr/local/lib/node_modules/jshint/bin/jshint --config jshintrc foo.js
foo.js: line 5, col 17, 'num' is not defined.
```

```
1 error
```

**Invariant:**

“If code passes JSHint, all top-level symbols resolve.”

**Strengthening the Invariant.** Can we do better? How about adding a pre-commit hook?

“If code is checked-in and commit hook ran,  
all top-level symbols resolve.”

Of course, sometimes the commit hook didn’t run. Better yet:

- Block deploys on test failures.

**Better invariant.**

“If code is deployed,  
all top-level symbols resolve.”

**Even better yet.** It is hard to tell whether code is deployed or not. Use git feature branches, merge when deployed.

“If code is in master,  
all top-level symbols resolve.”